

As you know, I have a 4090 video card. My goal is to create a 1 hour ambient music video 2x per week. Most of these types of videos on Youtube just have simple images. Some just use After Effects rain or snow effect over a still image. With new AI tools, some have motion such as "fly through" of a scene. I have used Grok Imagine to create some very cool video, but I need to upgrade to get more video creation output. I have not used ComfyUI in a while - can anything new with ComfyUI even come close to Higgsfield.ai or Sora high-end video? Should I even try? I have an idea to use After Effects and create cool images but then separate them into layers and create a 2.5d video clips for my ambient videos. As you know, with 2.5d video it is easy in AE to move the different layers at different speeds. But this takes time to edit. What would you suggest as a good workflow as I probably need to spend 1 day creating the music and 2 days creating the images, videos, and editing?

Based on your setup with a high-end 4090 GPU, aiming for two 1-hour ambient music videos per week is ambitious but feasible if we streamline the workflow around story-driven content. I'll frame this as a managed project process, incorporating AI agents for story arcs (focusing on emotional progression and user immersion) and screenplay (detailing scenes, text overlays, and motion cues). This aligns with creating videos that blend motion graphics, evolving visuals, and narrative text to guide the viewer through a cohesive tale—think subtle builds from calm introspection to ethereal climaxes, enhanced by your ambient music.

First, addressing ComfyUI for video generation: Recent developments (as of late 2025) have integrated strong open-source models like LTX-Video, Mochi 1, Pyramid Flow, CogVideoX-5B, Wan 2.2, and HunyuanVideo via custom nodes and extensions. These allow text-to-video, image-to-video, and even audio-driven generation, with workflows for high-res outputs up to 1080p at 8-30 seconds per clip. Your 4090 can handle these locally and efficiently—expect 5-10x faster inference than on weaker hardware, with options for batch processing and fine-tuning via extensions like VideoX-Fun for custom training.

However, these don't fully match the photorealism, coherent physics, or long-sequence consistency of closed high-end tools like OpenAI's Sora 2 or platforms like Higgsfield.ai (which aggregates top models including Sora 2, Veo 3.1, and WAN 2.5 for cinematic effects). ComfyUI's outputs can feel more artifact-prone or stylized, especially for complex "fly-through" scenes, based on community comparisons where users note Sora 2's superiority in natural motion and detail. That said, yes, you should try it—it's free, runs offline on your rig, and excels for iterative experimentation without quotas. Start with pre-built workflows for ambient styles (e.g., slow pans over landscapes) to test against your Grok Imagine videos. If quality falls short, use it as a supplement for quick prototypes rather than full clips.

For your 2.5D idea in After Effects (AE): It's a solid, time-tested approach for ambient videos—layering images and parallax motion creates depth without full 3D rendering. But as you noted, manual separation and keyframing eats time. To optimize, automate layering with AI tools like Photoshop's generative fill or ComfyUI's depth estimation nodes (via extensions like ControlNet for auto-masking), then import into AE for speed tweaks.

Suggested Workflow: 3-Day Cycle for One 1-Hour Video

This process treats each video as a mini-project: Start with story agents to ensure narrative flow, leverage AI for asset creation, and reserve editing for polish. Aim for 10-15 key scenes (4-6 minutes each) to fill the hour, syncing to your music's beats. Use Google Gemini or me (Grok) as writing agents—prompt one for arcs (e.g., "Craft a story arc for an ambient video about a forgotten forest: build from mystery to revelation, prioritizing serene UX") and the other for screenplay (e.g., "Write a scene-by-scene script with text overlays and motion descriptions"). Total target: 1 day music, 2 days visuals/editing, producing every 2-3 days to hit 2x/week.

Day 1: Music Creation (Full Day Focus)

- Compose or curate your 1-hour ambient track in your preferred DAW. Break it into sections (intro, build, peak, resolve) to align with the story arc.
- While finalizing, jot high-level story ideas based on the music's mood (e.g., oceanic journey for wavy synths).
- End by exporting stems or timestamps for syncing visuals later.

Day 2: Story Crafting, Image/Video Generation (Asset Build)

- **Morning: Story Agents (1-2 hours)** Use Gemini or me to generate the narrative backbone:
 - Agent 1 (Arc/UX): Define emotional progression—e.g., start with static calm, escalate to subtle motion, end in harmony. Focus on viewer immersion (e.g., text fades that evoke reflection).
 - Agent 2 (Screenplay): Output a 10-15 scene script with details like "Scene 3: Text 'Whispers of the Wind' overlays a layered forest image; foreground leaves drift slowly right, background static for depth." This ensures the video tells a story, not just loops visuals.
- **Midday: Generate Key Assets (3-4 hours)**
 - Use Grok Imagine (or upgrade to SuperGrok/Premium+ for higher limits) for base images—prompt for ambient scenes like "ethereal snowy landscape in cyberpunk style, high-res, layer-friendly." Generate 20-30 variations.
 - For motion: Test ComfyUI with models like Pyramid Flow for short "fly-through" clips (5-10s each) from your images. Install extensions via ComfyUI-Manager for Wan 2.2 or LTX-Video. Prompt: "Animate this image as a slow aerial fly-through with falling snow." Your 4090 will churn these out quickly (minutes per clip). If underwhelmed, fall back to static images.
 - Layer prep: In Photoshop or ComfyUI, auto-separate elements (e.g., sky, midground, foreground) using AI masks. Export as PSDs for AE.
- **Afternoon: Initial Motion Tests (2-3 hours)** Import layers to AE and create 2.5D prototypes: Set camera moves with different layer speeds (e.g., background at 0.5x). Add simple effects like rain/snow particles. Render low-res previews to check flow against music timestamps.

Day 3: Video Editing and Final Polish (Assembly)

- **Morning: Compile Clips (3-4 hours)** In Kdenlive (for quick timeline work) or AE, sequence your 2.5D clips, ComfyUI videos, and stills. Sync to music—use dissolves for ambient transitions, add motion graphics (e.g., AE's text animator for story overlays).
- **Midday: Effects and Text Integration (2-3 hours)** Layer in narrative text from the screenplay (e.g., subtle fades). Enhance with AE presets for glows, color grading, or particle effects over stills. If using ComfyUI clips, upscale them here for consistency.
- **Afternoon: Render and Review (1-2 hours)** Export the full 1-hour video at 1080p/60fps. Review for story coherence and UX—does it immerse without overwhelming? Tweak pacing if needed.

Efficiency Tips for Scaling to 2x/Week

- **Batch Processing:** In ComfyUI, queue multiple generations overnight on your 4090 to build an asset library (e.g., reusable snow/rain overlays).
- **Time Savers:** Skip full 2.5D for every scene—mix with simple AE effects over stills for 50% of the video. If Grok Imagine quotas limit you, pivot more to ComfyUI for unlimited local runs.
- **Project Management:** Track via a simple spreadsheet: Column for story arc, assets needed, generation time. Reuse scripts/templates across videos for themes (e.g., nature cycles).
- **Upgrades to Consider:** If video output is a bottleneck, Higgsfield.ai could bridge to Sora-level quality via their app (test their free tier), but stick local for control. For more Grok Imagine capacity, check x.ai/grok subscriptions.

This workflow balances creativity with efficiency, ensuring each video has a compelling story while leveraging your tools. If you share a sample story prompt or music snippet, I can simulate the agents to kickstart one.

22 web pages

Tool	Free Tier	Paid Plans (Monthly)	Annual Discounts (Standard)	Black Friday Deals (2025 Projections/Configured)	Key Fit for Your Workflow
Invideo.io	Limited: 4 exports/wk, watermark, 2 video mins/wk	Plus/Business: \$20 (basic edits, templates) Max/Unlimited: \$48+ (unlimited gens, stock media)	50% off annual (e.g., Business ~\$120/yr → ~\$36/yr; some sources note no BF, but promo codes for 30-40% off)	Up to 70% off annual plans (Nov 21-Dec 2), e.g., Business \$120/yr → ~\$36/yr; some sources note no BF, but promo codes for 30-40% off	Quick templates for ambient overlays/text; batch effects like rain over stills, reducing AE time. Test for 2.5D prototypes.
Grok (xAI)	Grok 3: Limited quotas	SuperGrok: \$30 (higher limits, Grok 4)	16-17% off (e.g., \$300/yr)	Projected 40% off annual (e.g., SuperGrok \$300/yr)	Boost Imagine output for story visuals; use as

Tool	Free Tier	Paid Plans (Monthly)	Annual Discounts (Standard)	Black Friday Deals (2025 Projections/Confirmed)	Key Fit for Your Workflow
Higgsfield.ai	Limited credits (e.g., 10-20 short clips/mo)	SuperGrok Heavy: ~\$40 (pro access) API: Token-based (e.g., \$0.30/M tokens for Grok 3 Mini)		→ \$180/yr, Nov 28-Dec 1); check https://x.ai/grok for details	screenplay agent in prompts. No video gen yet, but API for custom integrations.
OpenAI Sora 2	Invite-only: 30 gens/day (free, limited compute)	Pro: \$29 (600 credits, ~10-12 videos) Ultimate: \$49 (unlimited, faster gens)	40% off (e.g., Pro \$29/mo → \$17.40/mo, Nov 21-Dec 2); "biggest price drop in GenAI history"	Up to 65% off + 1yr unlimited images (e.g., Pro \$29/mo → \$17.40/mo, Nov 21-Dec 2); "biggest price drop in GenAI history"	Cinematic "fly-throughs" with Sora/Veo/WAN models; layer outputs into AE for 2.5D depth. Batch overnight on cloud.
Google AI Studio (Gemini)	Free: Unlimited basic use	Extra credits: \$4/10 gens Sora 2 Pro: \$200 (via ChatGPT Pro, up to 25s clips) API: \$0.10-0.50/sec Gemini Advanced: \$19.99 (via Google One AI Premium) Business: Per user/mo add-on Standard: ~\$15 (credits for clips) Pro: \$95 (unlimited, 4K) Enterprise: Custom	N/A (pay-per-use focus)	Nano Banana event: Up to 65% off bundles w/Sora access (e.g., via platforms like Mintly); no direct BF confirmed	Physics-realistic scenes for ambient builds; sync to music timestamps. Use Pro for watermark-free polish in Kdenlive.
RunwayML	Basic: Limited gens	Pro: \$35 (2000 credits, ~100 clips) Unlimited: \$95 (fast gens)	40-50% off annual	No confirmed BF; watch for promo codes or auto-discounts (e.g., 40% off annual projected)	Arc/UX agent for stories; Veo integration for motion clips. Free for scripting, paid for advanced gens.
Pika Labs	Basic: Limited credits	Pro: ~\$29 (600 credits)	N/A	Up to 50% off annual via bundles (e.g., AppSumo); check for flash deals	Editor-style for motion graphics; Gen-3/4 for parallax effects. Commercial rights for pro use.
Luma AI (Dream)	Free: Basic gens	Pro: ~\$29 (600 credits)	40% off annual	Up to 65% off in Nano Banana bundles; promo codes for 30-40%	Quick stylized shorts; 1.5 model for ambient fly-throughs. Roll-over credits for batching.
				Up to 50% off via ecosystem deals	Smooth motion for 3D-like scenes;

Tool	Free Tier	Paid Plans (Monthly)	Annual Discounts (Standard)	Black Friday Deals (2025 Projections/Configured)	Key Fit for Your Workflow
Machine)		Enterprise: Custom		(e.g., AppSumo)	depth for 2.5D layers in AE.

Updated Workflow Integration

- **Day 1 (Music):** Finalize track; use Gemini (free) as arc agent for mood-aligned story ideas.
- **Day 2 (Assets):** Grok as screenplay agent for scene scripts. Generate images via Grok Imagine/Higgsfield (BF deal for unlimited). For videos, mix ComfyUI (local, free) with Higgsfield/Sora 2 (cloud, credits for high-end fly-throughs). Batch 10-15 clips, auto-separate layers for AE import.
- **Day 3 (Editing):** Assemble in Kdenlive/AE with text from screenplay. If needed, Invideo for quick effects overlays.

Prioritize Higgsfield BF (65% off) for quota relief—test vs. ComfyUI for quality. Track costs in your project spreadsheet; aim for \$20-50/mo total.